

Applied Machine Learning

Lecture 1: Introduction



UNIVERSITY OF
GOTHENBURG

CHALMERS

Richard Johansson

January 16, 2018

why the fuss?

- ▶ media exposure; some impressive recent results
- ▶ snowball effect: everyone wants to do ML
- ▶ more data available
- ▶ lower barriers to entry: ML software is becoming user-friendly
- ▶ ML is more efficient because of improvements in hardware

topics covered in the course

- ▶ the usual “zoo”: a selection of machine learning models
 - ▶ what's the idea behind them?
 - ▶ how are they implemented? (at least on a high level)
 - ▶ what are the use cases?
 - ▶ how can we apply them practically?
- ▶ but hopefully also the “real-world context”:
 - ▶ extended “messy” practical assignments requiring that you think of what you're doing
 - ▶ (probably) 2 invited talks from industry
 - ▶ ethical and legal issues, interpretability

structure of teaching

- ▶ video lectures: mainly for theory
 - ▶ please watch the videos before each exercise session!
- ▶ lecture / exercise sessions (Tuesdays and Fridays)
 - ▶ some theory and introduction to ML software
 - ▶ interactive coding
 - ▶ solving exercises in groups
 - ▶ (tentatively) two industrial guest lectures
- ▶ lab sessions: you work on your assignments
 - ▶ please go to the 13-15 or the 15-17 session

assignments

- ▶ warmup exercise: quick tour of the scikit-learn library
- ▶ four compulsory **assignments**:
 1. “mini-project” where you solve a supervised learning task
 2. implement a classification algorithm
 3. neural network design
 4. written essay on ethics in ML
- ▶ please refer to the course PM for details about grading
- ▶ we will use the **Python** programming language
 - ▶ please ask for permission if you prefer to use something else

literature

- ▶ the main course book is *A Course in Machine Learning* by Hal Daumé III: <http://ciml.info>
- ▶ and additional papers to read for some topics
- ▶ example code will be posted on the course page

written exam on March 15

- ▶ a first part about basic concepts: you need to answer most of these questions correctly to pass
- ▶ a second part that requires more insight: answer these questions for a higher grade

why machine learning?

why would we want to “learn” the function from data instead of just implementing it?

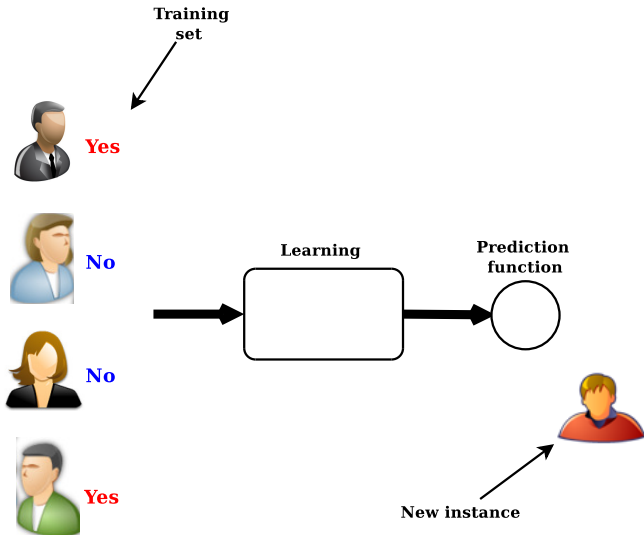
- ▶ usually **because we don't really know** how to write down the function by hand
 - ▶ speech recognition
 - ▶ image classification
 - ▶ machine translation
 - ▶ ...
- ▶ might not be necessary for **limited** tasks where we **know**
- ▶ what is more expensive in your case? knowledge or data?

don't forget your domain expertise!

machine learning automatizes some tasks, but we still need our brains:

- ▶ **defining** the tasks, terminology, evaluation metrics
- ▶ **annotating** training and testing data
- ▶ having an intuition about which **features** may be useful can be crucial
 - ▶ in general, features are more important than the choice of learning algorithm
- ▶ **error analysis**
- ▶ defining **constraints** to guide the learner


learning from data



example: is the patient diabetic?




example: is the patient diabetic?



- ▶ in order to predict, we make some measurements of properties we believe will be useful
 - ▶ these are called the **features**

decision tree classifiers

- ▶ a **decision tree** is a tree where
 - ▶ the internal nodes represent how we choose based on a feature
 - ▶ the leaves represent the return value of the classifier
- ▶ like the example we had previously:
 - ▶ IF the blood glucose level > 150 , THEN
 - ▶ IF the age > 50 , THEN return True
 - ▶ ELSE ...
 - ▶ ...




greedy decision tree learning (pseudocode)

```
def TrainDecisionTree( $T$ )  
  if  $T$  is unambiguous  
    return a leaf with the class of the examples in  $T$   
  if  $T$  has no features  
    return a leaf with the majority class of  $T$   
   $F \leftarrow$  the “most useful feature” in  $T$   
  for each possible value  $f_i$  of  $F$   
     $T_i \leftarrow$  the subset of  $T$  where  $F = f_i$   
    remove  $F$  from  $T_i$   
     $\text{tree}_i \leftarrow \text{TrainDecisionTree}(T_i)$   
  return a tree node that splits on  $F$ ,  
    where  $f_i$  is connected to the subtree  $\text{tree}_i$ 
```


problems with the naive approach

- ▶ ID3 and similar decision tree learning algorithms often have troubles with large, noisy datasets
- ▶ often, performance decreases with training set size!



- ▶ can be improved by using a separate development set:
 - ▶ **prune** the tree by removing the nodes that don't seem to matter for accuracy on the development set

scikit-learn toy example: a simple training set

```
# training set: the features
X = [{'city':'Gothenburg', 'month':'July'},
      {'city':'Gothenburg', 'month':'December'},
      {'city':'Paris', 'month':'July'},
      {'city':'Paris', 'month':'December'}]

# training set: the gold-standard outputs
Y = ['rain', 'rain', 'sun', 'rain']
```

scikit-learn toy example: training a classifier

```
from sklearn.feature_extraction import DictVectorizer
from sklearn.svm import LinearSVC
from sklearn.pipeline import make_pipeline
import pickle

pipeline = make_pipeline( DictVectorizer(), LinearSVC() )

# train the classifier
pipeline.fit(X, Y)

# optionally: save the classifier to a file...
with open('weather.classifier', 'wb') as f:
    pickle.dump(pipeline, f)
```


toy example: making new predictions and evaluating

```
from sklearn.metrics import accuracy_score

Xtest = [{'city':'Gothenburg', 'month':'June'},
          {'city':'Gothenburg', 'month':'November'},
          {'city':'Paris', 'month':'June'},
          {'city':'Paris', 'month':'November'}]

Ytest = ['rain', 'rain', 'sun', 'rain']

# classify all the test instances
guesses = pipeline.predict(Xtest)

# compute the classification accuracy
print(accuracy_score(Ytest, guesses))
```


some other practical functions

- ▶ making a training/test split:

```
from sklearn.cross_validation import train_test_split

train_files, dev_files = train_test_split(td_files,
                                         train_size=0.8,
                                         random_state=0)
```

- ▶ evaluation, e.g. accuracy, precision, recall, F-score:

```
from sklearn.metrics import f1_score

print(f1_score(Y_eval, Y_out))
```

- ▶ cross-validation over the training set:

```
from sklearn.cross_validation import cross_validate

cv_results = cross_validate(pipeline, X, Y)
```


types of machine learning problems: what are we predicting?

- ▶ **classification**: learning to output a category label
 - ▶ spam/non-spam; positive/negative; ...
- ▶ **regression**: learning to guess a number
 - ▶ value of a share; number of stars in a review; ...
- ▶ **structured prediction**: learning to build some structure
 - ▶ speech segmentation; machine translation; ...
- ▶ **ranking**: learn to order a set of items
 - ▶ search engines
- ▶ **reinforcement learning**: learning to act in an environment
 - ▶ dialogue systems; playing games; autonomous vehicles; ...

types of supervision (1): supervised



- ▶ in supervised learning, we have a labeled training set consists of input-output pairs
- ▶ our goal is to learn to imitate this labeling

age	workclass	lnwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	target
39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

	F1	F2	F3	F4	F5	F6	F7	F8	F9	RMSD
0	13558.30	4305.35	0.31754	162.1730	1.872791e+06	215.3590	4287.87	102	27.0302	17.284
1	6191.96	1623.16	0.26213	53.3894	8.034467e+05	87.2024	3328.91	39	38.5468	6.021
2	7725.98	1726.28	0.22343	67.2887	1.075648e+06	81.7913	2981.04	29	38.8119	9.275
3	8424.58	2368.25	0.28111	67.8325	1.210472e+06	109.4390	3248.22	70	39.0651	15.851
4	7460.84	1736.94	0.23280	52.4123	1.021020e+06	94.5234	2814.42	41	39.9147	7.962

types of supervision (2): **unsupervised**

- ▶ in **unsupervised** learning, we are given a set of “unorganized” data
- ▶ our goal is to discover some structure in the data





types of supervision (3): variations...

- ▶ **semisupervised** learning:
 - ▶ a small set of labeled examples plus a larger unlabeled set
- ▶ **active** learning:
 - ▶ the learning algorithm can ask for additional labeling of targeted examples
- ▶ **multitask** learning:
 - ▶ learning from closely related tasks

representation of the prediction function


we may represent our prediction function in different ways:

- ▶ **numerical** models:
 - ▶ weight vectors, probability tables
 - ▶ networked models
- ▶ **rule-based** models:
 - ▶ decision trees
 - ▶ rules expressed using logic



what goes on when we “learn”?

- ▶ the learning algorithm observes the examples in the training set
- ▶ it tries to find common patterns that explain the data: it **generalizes** so that we can make predictions for new examples




- ▶ how this is done depends on what algorithm we are using

why would we prefer “simple” hypotheses?



“overfitting” and “underfitting”: the bias–variance tradeoff



[Source: [Wikipedia](#)]

up next

- ▶ Thursday: lab session for the noncompulsory exercise
- ▶ topic of Friday's discussion: linear classifiers and regressors
- ▶ please prepare by watching the videos